PLSC 503 – Spring 2020
Nominal Outcomes

April 14, 2020

Motivation: Discrete *Outcomes*

$$\Pr(Y_i = j) = P_{ij}$$

$$\sum_{j=1}^{J} P_{ij} = 1$$

$$P_{ij} = \exp(\mathbf{X}_i \boldsymbol{\beta}_j)$$

Motivation, continued

Rescale:

$$\Pr(Y_i = j) \equiv P_{ij} = \frac{\exp(\mathbf{X}_i \boldsymbol{\beta}_j)}{\sum_{j=1}^{J} \exp(\mathbf{X}_i \boldsymbol{\beta}_j)}$$

Ensures

- $\Pr(Y_i = j) \in (0, 1)$
- $\sum_{j=1}^{J} \Pr(Y_i = j) = 1.0$

Constrain $\boldsymbol{\beta}_1 = \mathbf{0}$; then:

$$\Pr(Y_i = 1) = \frac{1}{1 + \sum_{j=2}^{J} \exp(\mathbf{X}_i \boldsymbol{\beta}_j')}$$

$$\Pr(Y_i = j) = \frac{\exp(\mathbf{X}_i \boldsymbol{\beta}_j')}{1 + \sum_{j=2}^{J} \exp(\mathbf{X}_i \boldsymbol{\beta}_j')}$$

where $\boldsymbol{\beta}_j' = \boldsymbol{\beta}_j - \boldsymbol{\beta}_1$.

Alternative Motivation: Discrete *Choice*

$$U_{ij} = \mu_i + \epsilon_{ij}$$

$$\mu_i = \mathbf{X}_i \boldsymbol{\beta}_j$$

$$
\begin{aligned}
\Pr(Y_i = j) &= \Pr(U_{ij} > U_{i\ell} \, \forall \, \ell \neq j \in J) \\
&= \Pr(\mu_i + \epsilon_{ij} > \mu_i + \epsilon_{i\ell} \, \forall \, \ell \neq j \in J) \\
&= \Pr(\mathbf{X}_i \boldsymbol{\beta}_j + \epsilon_{ij} > \mathbf{X}_i \boldsymbol{\beta}_\ell + \epsilon_{i\ell} \, \forall \, \ell \neq j \in J) \\
&= \Pr(\epsilon_{ij} - \epsilon_{i\ell} > \mathbf{X}_i \boldsymbol{\beta}_\ell - \mathbf{X}_i \boldsymbol{\beta}_j \, \forall \, \ell \neq j \in J)
\end{aligned}
$$

Discrete Choice (continued)

$\epsilon \sim$???

- *Type I Extreme Value*
- Density: $f(\epsilon) = \exp[-\epsilon - \exp(-\epsilon)]$
- CDF: $\int f(\epsilon) \equiv F(\epsilon) = \exp[-\exp(-\epsilon)]$

$$
\begin{aligned}
\Pr(Y_i = j) &= \Pr(U_j > U_1, U_j > U_2, ...U_j > U_J) \\
&= \int f(\epsilon_j) \left[\int_{-\infty}^{\epsilon_{ij}+\mathbf{X}_i\boldsymbol{\beta}_j-\mathbf{X}_i\boldsymbol{\beta}_1} f(\epsilon_1)d\epsilon_1 \times \int_{-\infty}^{\epsilon_{ij}+\mathbf{X}_i\boldsymbol{\beta}_j-\mathbf{X}_i\boldsymbol{\beta}_2} f(\epsilon_2)d\epsilon_2 \times ... \right] d\epsilon_j \\
&= \int f(\epsilon_j) \times \exp[-\exp(\epsilon_{ij} + \mathbf{X}_i\boldsymbol{\beta}_j - \mathbf{X}_i\boldsymbol{\beta}_1)] \times \\
&\quad \exp[-\exp(\epsilon_{ij} + \mathbf{X}_i\boldsymbol{\beta}_j - \mathbf{X}_i\boldsymbol{\beta}_2)] \times ...d\epsilon_j \\
&= \frac{\exp(\mathbf{X}_i\boldsymbol{\beta}_j)}{\sum_{j=1}^{J} \exp(\mathbf{X}_i\boldsymbol{\beta}_j)}
\end{aligned}
$$

Estimation

Define:

$$
\begin{aligned}
\delta_{ij} &= 1 \text{ if } Y_i = j, \\
&= 0 \text{ otherwise.}
\end{aligned}
$$

Then:

$$
\begin{aligned}
L_i &= \prod_{j=1}^{J} [\Pr(Y_i = j)]^{\delta_{ij}} \\
&= \prod_{j=1}^{J} \left[\frac{\exp(\mathbf{X}_i \boldsymbol{\beta}_j)}{\sum_{j=1}^{J} \exp(\mathbf{X}_i \boldsymbol{\beta}_j)} \right]^{\delta_{ij}}
\end{aligned}
$$

More Estimation

So:
$$L = \prod_{i=1}^{N} \prod_{j=1}^{J} \left[\frac{\exp(\mathbf{X}_i \boldsymbol{\beta}_j)}{\sum_{j=1}^{J} \exp(\mathbf{X}_i \boldsymbol{\beta}_j)} \right]^{\delta_{ij}}$$

and (of course):
$$\ln L = \sum_{i=1}^{N} \sum_{j=1}^{J} \delta_{ij} \ln \left[\frac{\exp(\mathbf{X}_i \boldsymbol{\beta}_j)}{\sum_{j=1}^{J} \exp(\mathbf{X}_i \boldsymbol{\beta}_j)} \right]$$

A (Descriptive) Example: 1992 Election

- 1992 National Election Study

- $Y \in \{\text{Bush} = 1, \text{Clinton} = 2, \text{Perot} = 3\}$

- $N = 1473$.

- $X =$ Party ID: $\{$ "Strong Democrats" $= 1 \rightarrow$ "Strong Republicans" $= 7\}$

MNL: 1992 Election ("Baseline" = Perot)

```
> nes92.mlogit<-vglm(presvote~partyid, multinomial, nes92)
> summary(nes92.mlogit)

Call:
vglm(formula = presvote ~ partyid, family = multinomial, data = nes92)

Coefficients:
              Estimate Std. Error z value          Pr(>|z|)
(Intercept):1  -1.8152     0.2456   -7.39   0.00000000000014 ***
(Intercept):2   3.0273     0.1783   16.98 < 0.0000000000000002 ***
partyid:1       0.4827     0.0476   10.15 < 0.0000000000000002 ***
partyid:2      -0.6805     0.0478  -14.25 < 0.0000000000000002 ***
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1   1

Number of linear predictors:  2
Names of linear predictors: log(mu[,1]/mu[,3]), log(mu[,2]/mu[,3])
Dispersion Parameter for multinomial family:    1
Residual deviance: 2167 on 2942 degrees of freedom
Log-likelihood: -1083 on 2942 degrees of freedom
Number of iterations: 5
```

MNL: 1992 Election ("Baseline" = Bush)

```
> Bush.nes92.mlogit<-vglm(formula = presvote~partyid,
        family=multinomial(refLevel=1),data=nes92)
> summary(Bush.nes92.mlogit)

Coefficients:
              Estimate Std. Error z value          Pr(>|z|)
(Intercept):1   4.8425     0.2373   20.41 < 0.0000000000000002 ***
(Intercept):2   1.8152     0.2456    7.39  0.00000000000000014 ***
partyid:1      -1.1632     0.0546  -21.32 < 0.0000000000000002 ***
partyid:2      -0.4827     0.0476  -10.15 < 0.0000000000000002 ***
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1   1

Number of linear predictors:  2
Names of linear predictors: log(mu[,2]/mu[,1]), log(mu[,3]/mu[,1])
Dispersion Parameter for multinomial family:    1
Residual deviance: 2167 on 2942 degrees of freedom
Log-likelihood: -1083 on 2942 degrees of freedom
Number of iterations: 5
```

MNL: 1992 Election ("Baseline" = Clinton)

```
> Clinton.nes92.mlogit<-vglm(formula=presvote~partyid,
                  family=multinomial(refLevel=2),data=nes92)
> summary(Clinton.nes92.mlogit)

Coefficients:
             Estimate Std. Error z value         Pr(>|z|)
(Intercept):1  -4.8425    0.2373   -20.4 <0.0000000000000002 ***
(Intercept):2  -3.0273    0.1783   -17.0 <0.0000000000000002 ***
partyid:1       1.1632    0.0546    21.3 <0.0000000000000002 ***
partyid:2       0.6805    0.0478    14.2 <0.0000000000000002 ***
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1   1

Number of linear predictors:  2
Names of linear predictors: log(mu[,1]/mu[,2]), log(mu[,3]/mu[,2])
Dispersion Parameter for multinomial family:    1
Residual deviance: 2167 on 2942 degrees of freedom
Log-likelihood: -1083 on 2942 degrees of freedom
Number of iterations: 5
```

Coefficient Estimates and "Baselines"

		"Baseline" category		
		Clinton	Perot	Bush
Comparison	Clinton	–	-0.68	-1.16
Category	Perot	0.68	–	-0.48
	Bush	1.16	0.48	–

It is exactly the same as the multinomial logit model. Period.

Choice-Specific Covariates

```
> library(mlogit)
> colnames(nes92)<-c("caseid","presvote","partyid","FT.Bush",
  "FT.Clinton","FT.Perot")
> nes92$PVote<-factor(nes92$presvote,labels=c("Bush","Clinton","Perot"))
> nes92CL<-mlogit.data(nes92,shape="wide",choice="PVote",varying=4:6)
> head(nes92)

  caseid presvote partyid FT.Bush FT.Clinton FT.Perot      NA   PVote
1   3001        1       6      85         30        0    Bush    Bush
2   3002        1       7     100          0        0    Bush    Bush
3   3003        1       7      85         30       60    Bush    Bush
4   3005        2       6      40         60       60 Clinton Clinton
5   3006        2       2      30         70       50 Clinton Clinton
6   3007        2       1      15         70       50 Clinton Clinton
```

$$\Pr(Y_{ij} = 1) = \frac{\exp(\mathbf{Z}_{ij}\gamma)}{\sum_{j=1}^{J} \exp(\mathbf{Z}_{ij}\gamma)}$$

Combinations: $\mathbf{X}_i\boldsymbol{\beta}$ and $\mathbf{Z}_{ij}\gamma$

- "Fixed effects"

- Observation-specific $\mathbf{X}$s

- Interactions…

```
> nes92.clogit<-mlogit(PVote~FT|partyid,data=nes92CL)
> summary(nes92.clogit)

Call:
mlogit(formula = PVote ~ FT | partyid, data = nes92CL, method = "nr",
    print.level = 0)

Frequencies of alternatives:
  Bush Clinton  Perot
 0.339   0.469  0.191

nr method
6 iterations, 0h:0m:0s
g'(-H)^-1g = 0.00293
successive function values within tolerance limits

Coefficients :
                   Estimate Std. Error t-value               Pr(>|t|)
Clinton:(intercept)  2.81272    0.26880   10.46 < 0.0000000000000002 ***
Perot:(intercept)    0.94353    0.28563    3.30              0.00096 ***
FT                   0.06299    0.00322   19.58 < 0.0000000000000002 ***
Clinton:partyid     -0.63187    0.06225  -10.15 < 0.0000000000000002 ***
Perot:partyid       -0.19212    0.05703   -3.37              0.00076 ***
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1   1

Log-Likelihood: -736
McFadden R^2:  0.519
Likelihood ratio test : chisq = 1590 (p.value = <0.0000000000000002)
```

Interpretation: Example Data Redux

- 1992 ANES ($N = 1473$)

- Variables:
  - presvote: 1=Bush, 2=Clinton, 3=Perot
  - partyid: (seven-point scale, 7=GOP)
  - age (in years)
  - white (naturally coded)
  - female (ditto)

Baseline MNL Results: 1992 Election

```
> NES.MNL<-vglm(presvote~partyid+age+white+female,data=BigNES92,
+               multinomial(refLevel=1))
> summaryvglm(NES.MNL)

Call:
vglm(formula = presvote ~ partyid + age + white + female, family = multinomial(refLevel = 1),
    data = BigNES92)

Coefficients:
               Estimate Std. Error z value          Pr(>|z|)
(Intercept):1   5.80665    0.44301   13.11 < 0.0000000000000002 ***
(Intercept):2   1.98008    0.52454    3.77           0.00016 ***
partyid:1      -1.13561    0.05486  -20.70 < 0.0000000000000002 ***
partyid:2      -0.50132    0.04870  -10.29 < 0.0000000000000002 ***
age:1          -0.00260    0.00514   -0.51           0.61276
age:2          -0.01556    0.00504   -3.09           0.00203 **
whiteWhite:1   -0.98908    0.31346   -3.16           0.00160 **
whiteWhite:2    0.87918    0.43605    2.02           0.04377 *
female:1       -0.12500    0.16895   -0.74           0.45936
female:2       -0.50928    0.16266   -3.13           0.00174 **
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1   1

Number of linear predictors: 2
Names of linear predictors: log(mu[,2]/mu[,1]), log(mu[,3]/mu[,1])
Dispersion Parameter for multinomial family:    1
Residual deviance: 2107 on 2936 degrees of freedom
Log-likelihood: -1054 on 2936 degrees of freedom
Number of iterations: 5
```

Global $\ln LR$ statistic Q tests:

$$\hat{\boldsymbol{\beta}} = \mathbf{0} \; \forall j, k$$

$$Q \sim \chi^2_{(J-1)(k-1)}$$

Test H: No Effect of age

```
> library(aod)
> wald.test(b=c(t(coef(NES.MNL))),Sigma=vcov(NES.MNL),Terms=c(5,6))

Wald test:
----------

Chi-squared test:
X2 = 11.0, df = 2, P(> X2) = 0.0042
```

Test H: No Difference – Clinton vs. Bush

```
> wald.test(b=c(t(coef(NES.MNL))),Sigma=vcov(NES.MNL),Terms=c(1,3,5,7,9))

Wald test:
----------

Chi-squared test:
X2 = 444.6, df = 5, P(> X2) = 0.0
```

Predicted <u>Outcomes</u>

```
> PickBush<-ifelse(fitted.values(NES.MNL)[,1]>fitted.values(NES.MNL)[,2]
        & fitted.values(NES.MNL)[,1]>fitted.values(NES.MNL)[,3], 1,0)
> PickWJC<-ifelse(fitted.values(NES.MNL)[,2]>fitted.values(NES.MNL)[,1]
        & fitted.values(NES.MNL)[,2]>fitted.values(NES.MNL)[,3], 2, 0)
> PickHRP<-ifelse(fitted.values(NES.MNL)[,3]>fitted.values(NES.MNL)[,1]
        & fitted.values(NES.MNL)[,3]>fitted.values(NES.MNL)[,2], 3, 0)

> OutHat<-PickBush+PickWJC+PickHRP
> table(BigNES92$presvote,OutHat)

  OutHat
     1   2   3
1 415  77   8
2  56 619  16
3 135 133  14
```

- "Null" Model: $\left(\frac{691}{1473}\right) = 46.9\%$ correct.

- Estimated model: $\frac{(415+619+14)}{1473} = \frac{1048}{1473} = 71.2\%$ correct.

- PRE $= \frac{1048-691}{1473-691} = \frac{357}{782} = 45.7\%$.

- Correct predictions: 90% Clinton, 83% Bush, 5% Perot.

Marginal Effects

$$\frac{\partial \Pr(Y_i = j)}{\partial X_k} = \Pr(Y_i = j|\mathbf{X}) \left[\hat{\boldsymbol{\beta}}_{jk} - \sum_{j=1}^{J} \hat{\boldsymbol{\beta}}_{jk} \times \Pr(Y_i = j|\mathbf{X}) \right]$$

Depends on:

- $\widehat{\Pr(Y_i = j)}$

- $\hat{\beta}_{jk}$

- $\sum_{j=1}^{J} \hat{\boldsymbol{\beta}}_{jk}$

See the end for (Stata) examples...

Odds ("Relative Risk") Ratios

$$\ln\left[\frac{\Pr(Y_i = j|\mathbf{X})}{\Pr(Y_i = j'|\mathbf{X})}\right] = \mathbf{X}(\hat{\boldsymbol{\beta}}_j - \hat{\boldsymbol{\beta}}_{j'})$$

Setting $\hat{\boldsymbol{\beta}}_{j'} = \mathbf{0}$:

$$\ln\left[\frac{\Pr(Y_i = j|\mathbf{X})}{\Pr(Y_i = j'|\mathbf{X})}\right] = \mathbf{X}\hat{\boldsymbol{\beta}}_j$$

One-Unit Change in X_k:

$$RRR_{jk} = \exp(\boldsymbol{\beta}_{jk})$$

δ-Unit Change in X_k:

$$RRR_{jk} = \exp(\boldsymbol{\beta}_{jk} \times \delta)$$

Odds ("Relative Risk") Ratios

```
> mnl.or <- function(model) {
    coeffs <- c(t(coef(model)))
    lci <- exp(coeffs - 1.96 * diag(vcov(NES.MNL))^0.5)
    or <- exp(coeffs)
    uci <- exp(coeffs + 1.96* diag(vcov(NES.MNL))^0.5)
    lreg.or <- cbind(lci, or, uci)
    lreg.or
  }

> mnl.or(NES.MNL)
                  lci      or      uci
(Intercept):1 139.5398 332.5036 792.3088
(Intercept):2   2.5909   7.2433  20.2504
partyid:1       0.2885   0.3212   0.3577
partyid:2       0.5506   0.6057   0.6664
age:1           0.9874   0.9974   1.0075
age:2           0.9749   0.9846   0.9943
whiteWhite:1    0.2012   0.3719   0.6875
whiteWhite:2    1.0248   2.4089   5.6623
female:1        0.6337   0.8825   1.2289
female:2        0.4369   0.6009   0.8266
```

Odds Ratios: Interpretation

- A one unit increase in `partyid` corresponds to:

  - A decrease in the odds of a Clinton vote, versus a vote for Bush, of $\exp(-1.136) = 0.321$ (or about 68 percent), and
  - A decrease in the odds of a Perot vote, versus a vote for Bush, of $\exp(-0.501) = 0.606$ (or about 40 percent).
  - These are *large* decreases in the odds – not surprisingly, more Republican voters are *much* more likely to vote for Bush than for Perot or Clinton.

- Similarly, `female` voters are:

  - No more or less likely to vote for Clinton vs. Bush (OR=0.88), but
  - Roughly 40 percent less likely to have voted for Perot (OR=0.60).
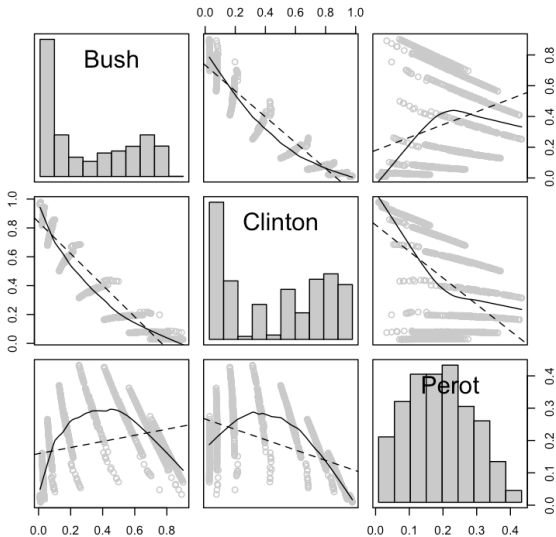
Predicted Probabilities

$$
\begin{aligned}
\Pr(\widehat{\mathtt{presvote}_i} = \mathsf{Bush}) &= \frac{\exp(\mathbf{X}_i \hat{\boldsymbol{\beta}}_{\mathsf{Bush}})}{\sum_{j=1}^{J} \exp(\mathbf{X}_i \hat{\boldsymbol{\beta}}_j)} \\
&= \frac{1}{1 + \sum_{j=2}^{J} \exp(\mathbf{X}_i \hat{\boldsymbol{\beta}}_j)}
\end{aligned}
$$

In-Sample Predicted Probabilities

```
> hats<-as.data.frame(fitted.values(NES.MNL))
> names(hats)[3]<-"Perot" # nice names...
> names(hats)[2]<-"Clinton"
> names(hats)[1]<-"Bush"
> attach(hats)

> library(car)
> scatterplot.matrix(~Bush+Clinton+Perot,
  diagonal="histogram",col=c("black","grey"))
```

```
> nes92.clogit<-mlogit(PVote~FT|partyid,data=nes92CL)
> summary(nes92.clogit)

Call:
mlogit(formula = PVote ~ FT | partyid, data = nes92CL, method = "nr",
    print.level = 0)

nr method
6 iterations, 0h:0m:0s
g'(-H)^-1g = 0.00293
successive function values within tolerance limits

Coefficients :
                    Estimate Std. Error t-value
Clinton:(intercept)  2.81272    0.26880   10.46
Perot:(intercept)    0.94353    0.28563    3.30
FT                   0.06299    0.00322   19.58
Clinton:partyid     -0.63187    0.06225  -10.15
Perot:partyid       -0.19212    0.05703   -3.37
                                 Pr(>|t|)
Clinton:(intercept) < 0.0000000000000002 ***
Perot:(intercept)                0.00096 ***
FT                  < 0.0000000000000002 ***
Clinton:partyid     < 0.0000000000000002 ***
Perot:partyid                    0.00076 ***
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1   1

Log-Likelihood: -736
McFadden R^2:  0.519
Likelihood ratio test : chisq = 1590 (p.value = <0.0000000000000002)
```

Predicted Probabilities (In-Sample)

```
> CLhats<-predict(NES.CL,type="expected")
> plot(cldata$FT,CLhats,xlab="Feeling Thermometer",ylab="Predicted Probability")
> lines(lowess(CLhats~cldata$FT),lwd=2,col="red")
```

Other Topics (for PLSC 504)

- "Independence of Irrelevant Alternatives"

- $\rightarrow$ Multinomial Probit

- $\rightarrow$ Heteroscedastic Extreme Value model

- "Mixed" Logit

- Nested Logit